

Big Data Project, Summer 2021

Course Title: Big Data

Course No.: PM-ASDS06

Submit to –

Dr. Md. Rezaul Karim

Associate Professor, Department of Statistics

Submitted by –

Md. Shamimul Islam ID:20204012

K. M. Sharif Ahmed ID: 20204062

Mirza Rahat ID: 20204036

Minhajur Rahman Khan ID: 20204042

Section-B, 4 Batch

Professional Masters in Applied Statistics and Data Science (ASDS) Jahangirnagar University

Contents:

1. Perceptron algorithm for single-layer neural network
2. Three-layer Artificial neural networks with SGD and Multi-Class Cross-Entropy Loss
3. Three-layer Artificial neural networks with SGD and Negative Log-Likelihood Loss Function
4. Three-layer Artificial neural networks with ADagrad and Multi-Class Cross-Entropy Loss
5. Three-layer Artificial neural networks with Adagrad and Negative Log-Likelihood Loss Function
6. Three-layer Artificial neural networks with RMSprop and Multi-Class Cross-Entropy Loss
7. Three-layer Artificial neural networks with RMSprop and Negative Log-Likelihood Loss Function
8. Three-layer Artificial neural networks with Adam and Multi-Class Cross-Entropy Loss
9. Three-layer Artificial neural networks with Adam and Negative Log-Likelihood Loss Function
10. Comparison Table for ANNs with different Loss functions and optimizers
11. CNNs with SGD and Multi-Class Cross-Entropy Loss
12. CNNs with SGD and Negative Log-Likelihood Loss Function
13. CNNs with ADagrad and Multi-Class Cross-Entropy Loss
14. CNNs with Adagrad and Negative Log-Likelihood Loss Function
15. CNNs with RMSprop and Multi-Class Cross-Entropy Loss
16. CNNs with RMSprop and Negative Log-Likelihood Loss Function
17. CNNs with Adam and Multi-Class Cross-Entropy Loss
18. CNNs with Adam and Negative Log-Likelihood Loss Function
19. Comparison Table for CNNs with different Loss function and optimizer

In []:

```
from google.colab import drive
```

```
drive.mount('/content/drive')
```

Mounted at /content/drive

In []:

```
import numpy as np
import gzip
import os
import sys
import struct
import numpy as np

def read_image(fi):
    magic, n, rows, columns = struct.unpack(">IIII", fi.read(16))
    assert magic == 0x00000803
    assert rows == 28
    assert columns == 28
    rawbuffer = fi.read()
    assert len(rawbuffer) == n * rows * columns
    rawdata = np.frombuffer(rawbuffer, dtype='>u1', count=n*rows*columns)
    return rawdata.reshape(n, rows, columns).astype(np.float32)

def read_label(fi):
    magic, n = struct.unpack(">II", fi.read(8))
    assert magic == 0x00000801
    rawbuffer = fi.read()
    assert len(rawbuffer) == n
    return np.frombuffer(rawbuffer, dtype='>u1', count=n)

train_x = '/content/drive/MyDrive/Dataset/mnist/train-images-idx3-ubyte.gz'
train_y = "/content/drive/MyDrive/Dataset/mnist/train-labels-idx1-ubyte.gz"
test_x = "/content/drive/MyDrive/Dataset/t10k-images-idx3-ubyte.gz"
test_y = "/content/drive/MyDrive/Dataset/t10k-labels-idx1-ubyte.gz"

np.savez_compressed(
    'mnist',
    train_x=read_image(gzip.open(train_x, 'rb')),
    train_y=read_label(gzip.open(train_y, 'rb')),
    test_x=read_image(gzip.open(test_x, 'rb')),
    test_y=read_label(gzip.open(test_y, 'rb'))
)
```

In []:

```
data = np.load('mnist.npz')

print(data['train_x'].shape, data['train_x'].dtype)
print(data['train_y'].shape, data['train_y'].dtype)
print(data['test_x'].shape, data['test_x'].dtype)
print(data['test_y'].shape, data['test_y'].dtype)
#print(data['train_x'])
#j = data['train_x']
#print(j[1])
```

```
(60000, 28, 28) float32
(60000,) uint8
(10000, 28, 28) float32
(10000,) uint8
```

Visualize the dataset

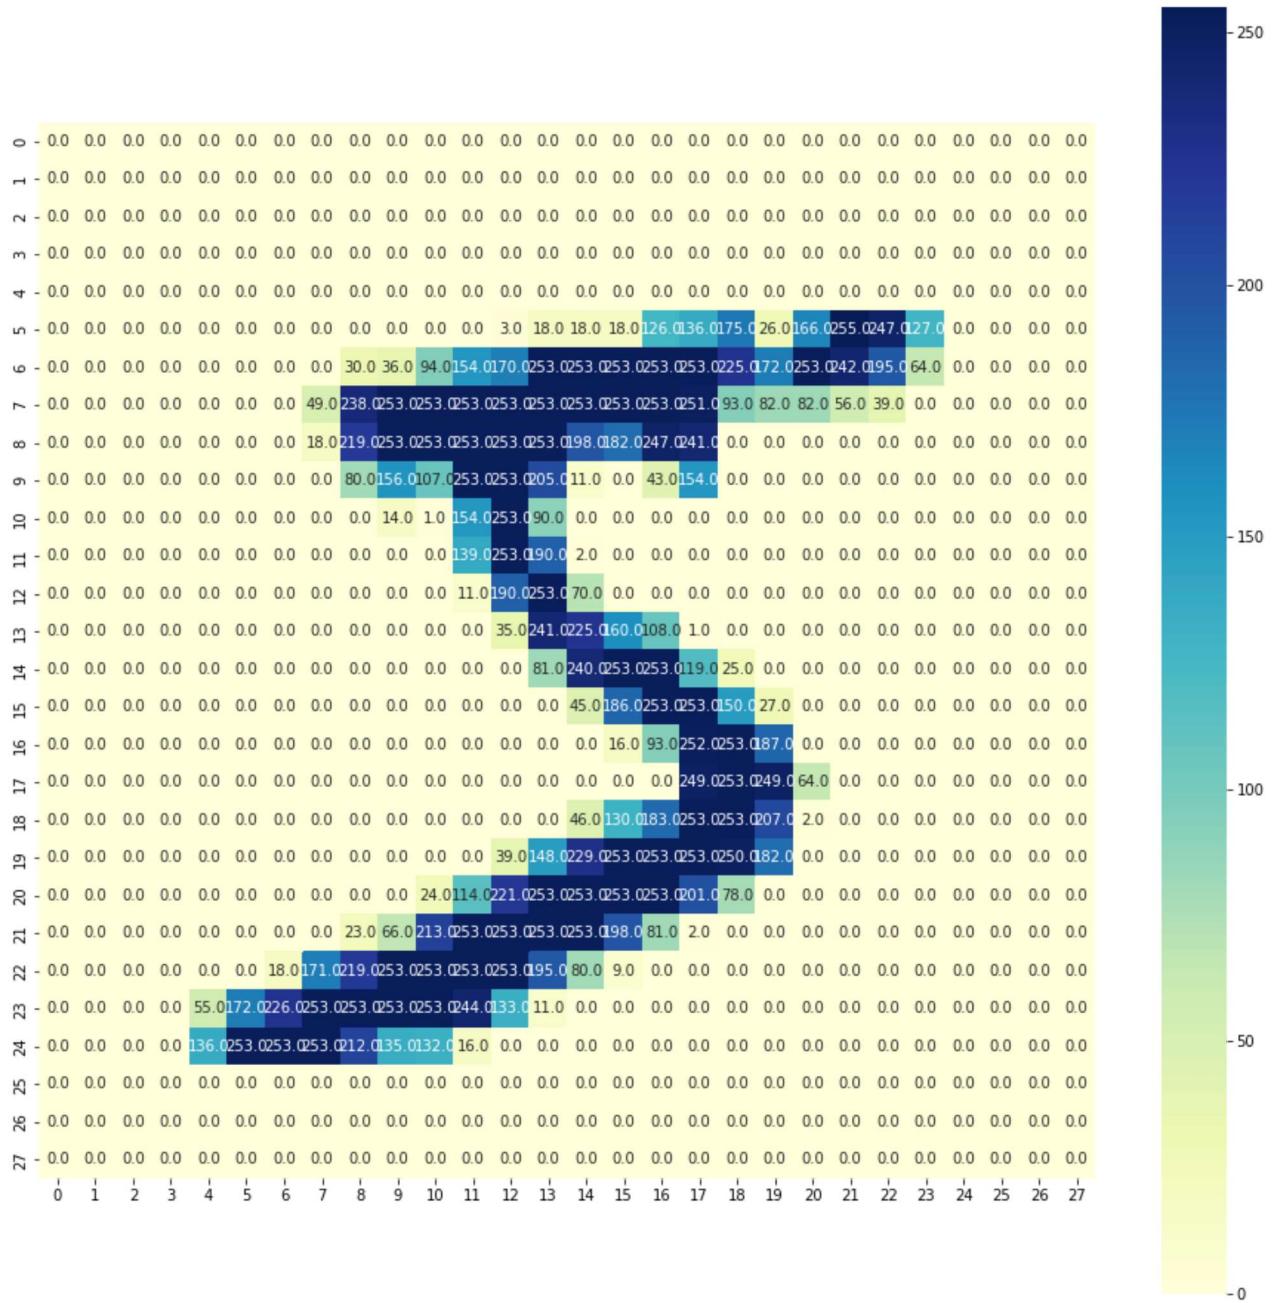
```
In [ ]: import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns

# Index number of an instance (change this to view another instance).
i = 0

data = np.load('mnist.npz')
image = data['train_x'][i]
label = data['train_y'][i]

print(label)
f, ax = plt.subplots(figsize=(16, 16))
sns.heatmap(image, annot=True, fmt='.1f', square=True, cmap="YlGnBu")
plt.show()
```

5



In []: `!pip install livelossplot`

Perceptron algorithm for single-layer neural network

```
In [ ]:
import numpy as np
from livelossplot import PlotLosses

def image_to_vector(X):
    X = np.reshape(X, (len(X), -1))      # Flatten: (N x 28 x 28) -> (N x 784)
    return np.c_[X, np.ones(len(X))]      # Append 1: (N x 784) -> (N x 785)

data = np.load('mnist.npz')
Xtrain = image_to_vector(data['train_x'])
Ytrain = data['train_y']
Xtest = image_to_vector(data['test_x'])
Ytest = data['test_y']
```

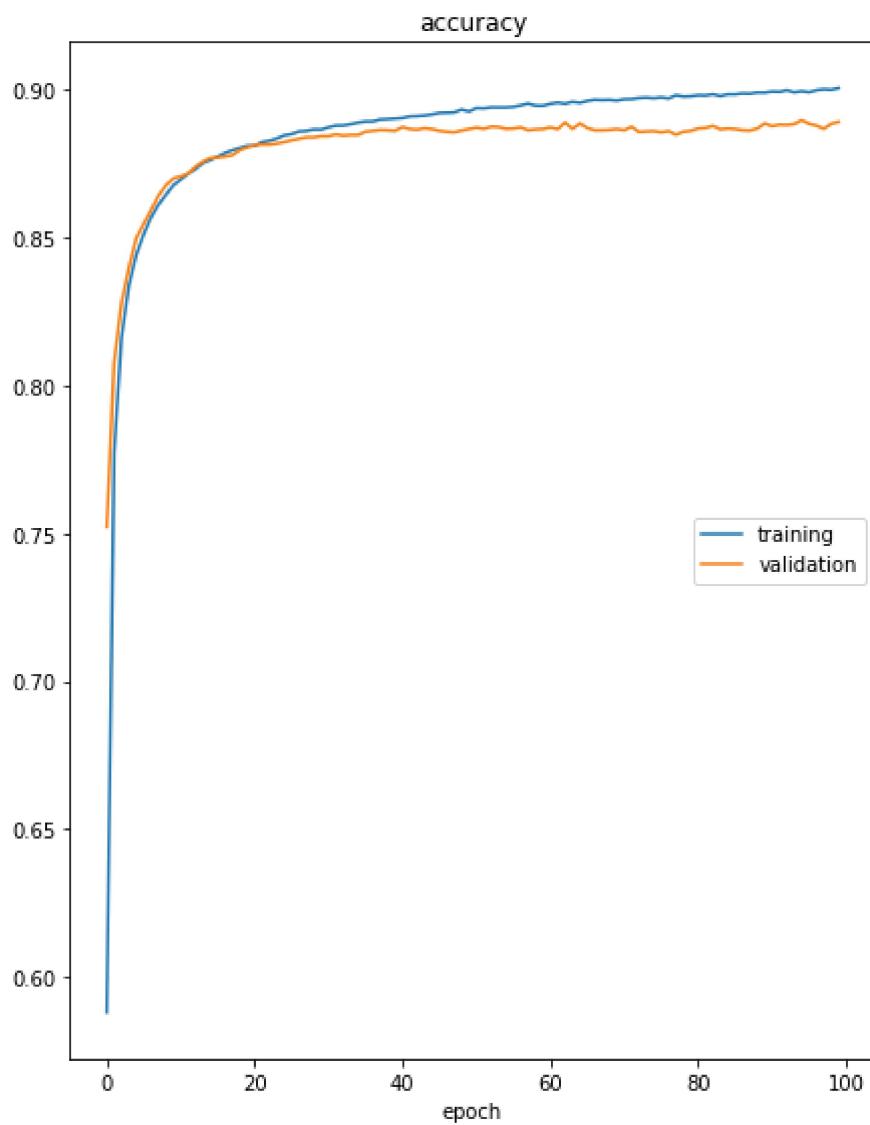
```
W = np.random.randn(10, 28*28+1)

eta = 0.001
liveloss = PlotLosses()
for t in range(100):
    # Structured perceptron for updating weights.
    num_correct_train = 0
    for x, y in zip(Xtrain, Ytrain):
        y_pred = np.argmax(np.dot(W, x))
        if y_pred != y:
            W[y] += x * eta
            W[y_pred] -= x * eta
        else:
            num_correct_train += 1

    # Evaluate and store the accuracy on the test set.
    num_correct_test = 0
    for x, y in zip(Xtest, Ytest):
        y_pred = np.argmax(np.dot(W, x))
        if y_pred == y:
            num_correct_test += 1

    # Visualize accuracy values on the training and test sets.
    liveloss.update({
        'accuracy': float(num_correct_train) / len(Ytrain),
        'val_accuracy': float(num_correct_test) / len(Ytest)
    })
    liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(
    float(num_correct_test) / len(Ytest),
    float(num_correct_train) / len(Ytrain)
))
```



```
accuracy
  training          (min: 0.588, max: 0.900, cur: 0.900)
  validation        (min: 0.752, max: 0.890, cur: 0.889)
Accuracy: 0.8890 (test), 0.9004 (train)
```

Stochastic gradient descent for single-layer neural network

In []:

```
import numpy as np
from livelossplot import PlotLosses

def softmax(x):
    # Result of softmax are invariant even if we add/subtract a constant.
    ex = np.exp(x - np.max(x)) # Subtract such that the maximum value is one.
    return ex / ex.sum(axis=0)

def image_to_vector(X):
    X = np.reshape(X, (len(X), -1))      # Flatten: (N x 28 x 28) -> (N x 784)
    return np.c_[X, np.ones(len(X))]      # Append 1: (N x 784) -> (N x 785)

def label_to_onehot(Y, K):
    return np.eye(K)[Y]                  # e.g., 3 -> [0, 0, 0, 1, 0, 0, 0, 0, 0]
```

```
data = np.load('mnist.npz')
Xtrain = image_to_vector(data['train_x'])
Ytrain = label_to_onehot(data['train_y'], 10)
Xtest = image_to_vector(data['test_x'])
Ytest = data['test_y']

W = np.random.randn(10, 28*28+1)

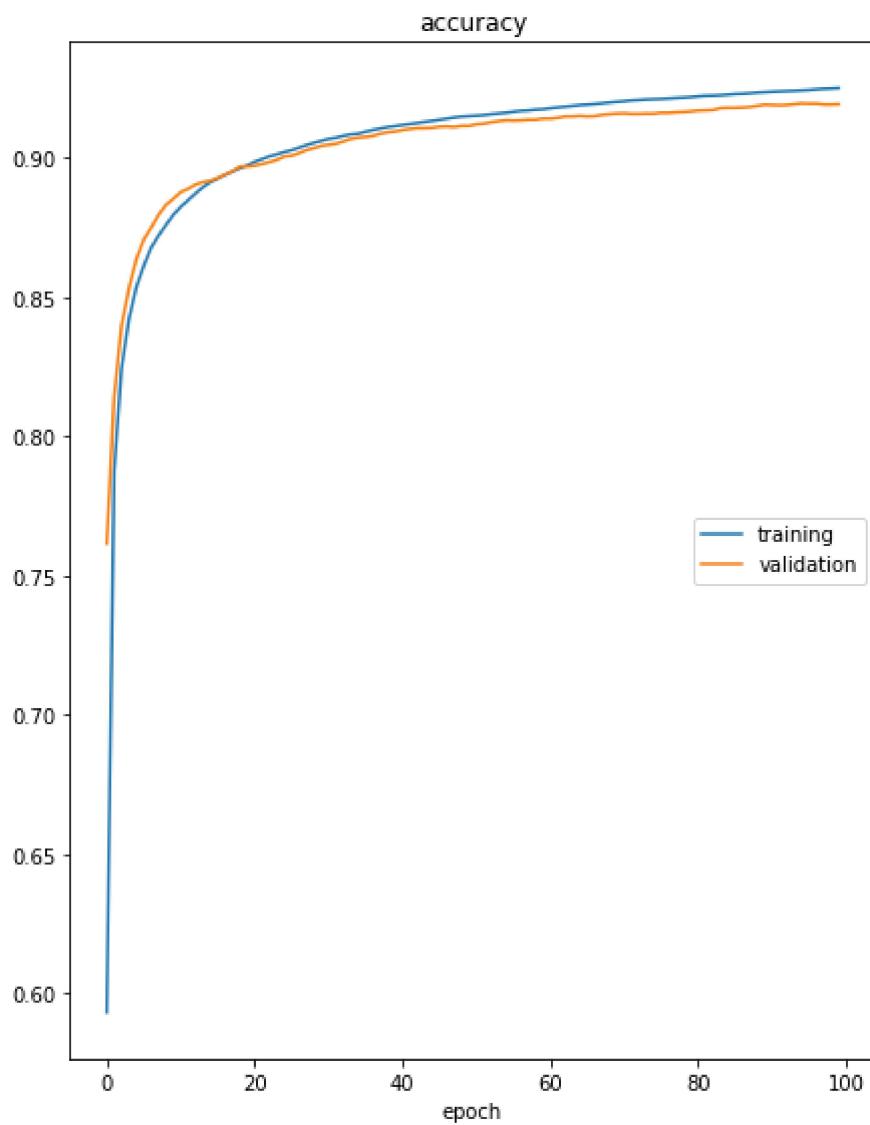
eta = 0.001
liveloss = PlotLosses()
for t in range(100):
    num_correct_train = 0

    # Stochastic gradient descent.
    for x, y in zip(Xtrain, Ytrain):
        y_pred = softmax(np.dot(W, x))
        W += np.outer(eta * (y - y_pred), x)
        if np.argmax(y_pred) == np.argmax(y):
            num_correct_train += 1

    # Evaluate and store the accuracy on the test set.
    num_correct_test = 0
    for x, y in zip(Xtest, Ytest):
        y_pred = np.argmax(np.dot(W, x))
        if y_pred == y:
            num_correct_test += 1

    # Visualize accuracy values on the training and test sets.
    liveloss.update({
        'accuracy': float(num_correct_train) / len(Ytrain),
        'val_accuracy': float(num_correct_test) / len(Ytest)
    })
    liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(
    float(num_correct_test) / len(Ytest),
    float(num_correct_train) / len(Ytrain)
))
```



```
accuracy
  training          (min: 0.593, max: 0.925, cur: 0.925)
  validation        (min: 0.761, max: 0.920, cur: 0.919)
Accuracy: 0.9193 (test), 0.9251 (train)
```

Training with pytorch

Convert the numpy arrays into pytorch tensor

```
In [ ]:
def create_dataset(x, y, flatten=False):
    if flatten:
        # Convert it into a matrix (N [samples], 28*28 [dims])
        xt = torch.from_numpy(x).view(len(x), -1)
    else:
        # Convert it into a 4D tensor (N [samples], 1 [ch], 28 [px], 28 [px])
        xt = torch.from_numpy(x).unsqueeze(1)
    yt = torch.from_numpy(y).long()
    return TensorDataset(xt, yt)
```

```

def test_model(model, loss_fn, test_loader, device):
    model.eval()

    loss = 0.
    num_correct = 0.
    for batch_idx, (x, y) in enumerate(test_loader):
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        loss += loss_fn(y_pred, y).item()
        _, predicted = torch.max(y_pred.data, 1)
        num_correct += (predicted == y).sum().item()

    model.train()
    loss /= len(test_loader.dataset)
    num_correct /= len(test_loader.dataset)
    return loss, num_correct

```

Single-layer neural network

In []:

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

def create_dataset(x, y):
    xt = torch.from_numpy(x).view(len(x), -1)
    yt = torch.from_numpy(y).long()
    return TensorDataset(xt, yt)

model = nn.Sequential()
model.add_module('fc1', nn.Linear(784, 10, bias=True))
model.to(device)

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'])
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)

loss_fn = nn.CrossEntropyLoss(size_average=False)
optimizer = optim.SGD(model.parameters(), lr=0.001)

for t in range(100):
    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)

        # Compute the loss value.
        loss = loss_fn(y_pred, y)

        # Update the parameters.
        optimizer.zero_grad()

```

```
loss.backward()
optimizer.step()
```

```
/usr/local/lib/python3.7/dist-packages/torch/nn/_reduction.py:42: UserWarning: size_average and reduce args will be deprecated, please use reduction='sum' instead.
  warnings.warn(warning.format(ret))
```

In []:

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

model = nn.Sequential()
model.add_module('fc1', nn.Linear(784, 10, bias=True))
print(model)
model.to(device)

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'], flatten=True)
test_set = create_dataset(data['test_x'], data['test_y'], flatten=True)
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.CrossEntropyLoss(size_average=False)
optimizer = optim.SGD(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

        # Compute the Loss value.
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

        # Update the parameters.
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    # Compute the average Loss and accuracy.
    train_loss /= len(train_loader.dataset)
    train_correct /= float(len(train_loader.dataset))

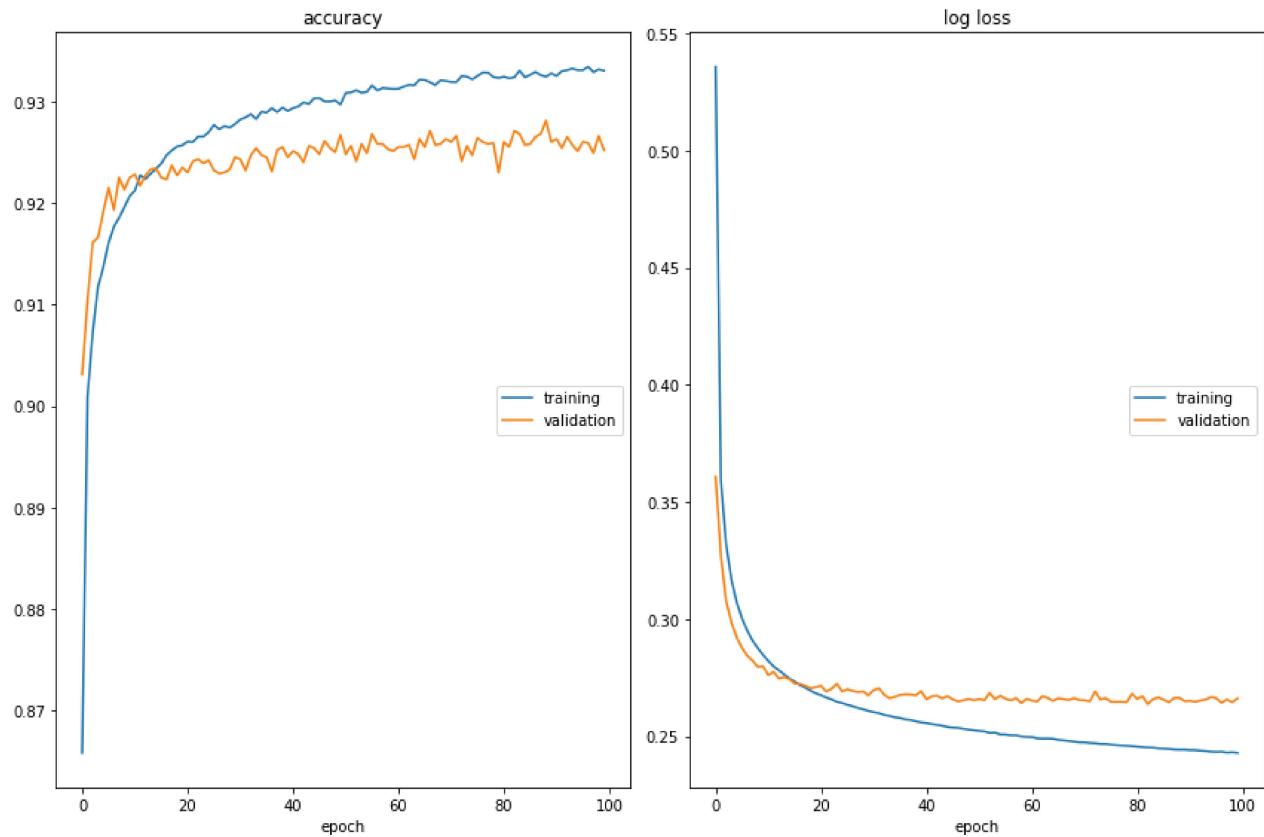
    # Evaluate the model on the test set.
    test_loss, test_correct = test_model(model, loss_fn, test_loader, device)
```

```

# Visualize the loss and accuracy values.
liveloss.update({
    'log loss': train_loss,
    'val_log loss': test_loss,
    'accuracy': train_correct,
    'val_accuracy': test_correct,
})
liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))

```



```

accuracy
    training          (min: 0.866, max: 0.933, cur: 0.933)
    validation        (min: 0.903, max: 0.928, cur: 0.925)
log loss
    training          (min: 0.243, max: 0.536, cur: 0.243)
    validation        (min: 0.264, max: 0.361, cur: 0.266)
Accuracy: 0.9252 (test), 0.9330 (train)

```

Three-layer neural network

For Stochastic Gradient Descent (SGD) and Multi-Class Cross-Entropy Loss

In []:

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader

```

```

import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

model = nn.Sequential()
model.add_module('fc1', nn.Linear(784, 256))
model.add_module('relu1', nn.ReLU())
model.add_module('dropout1', nn.Dropout())
model.add_module('fc2', nn.Linear(256, 256))
model.add_module('relu2', nn.ReLU())
model.add_module('dropout2', nn.Dropout())
model.add_module('fc3', nn.Linear(256, 10))
print(model)
model.to(device)

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'], flatten=True)
test_set = create_dataset(data['test_x'], data['test_y'], flatten=True)
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.CrossEntropyLoss(size_average=False)
optimizer = optim.SGD(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

        # Compute the Loss value.
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

        # Update the parameters.
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    # Compute the average Loss and accuracy.
    train_loss /= len(train_loader.dataset)
    train_correct /= float(len(train_loader.dataset))

    # Evaluate the model on the test set.
    test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

    # Visualize the Loss and accuracy values.
    liveloss.update({
        'log loss': train_loss,
        'val_log loss': test_loss,
        'accuracy': train_correct,
        'val_accuracy': test_correct,
    })

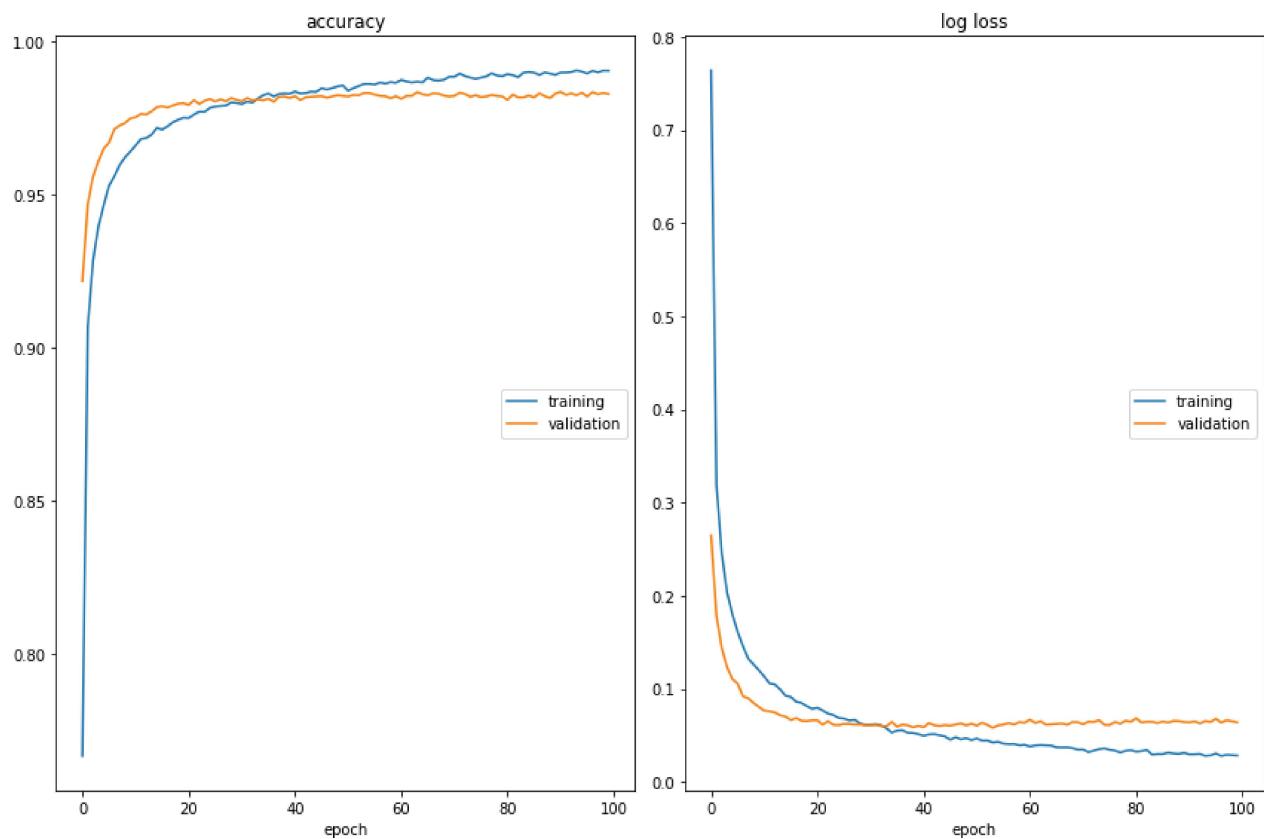
```

```

        })
    liveLoss.draw()

    print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))

```



```

accuracy
    training          (min: 0.767, max: 0.991, cur: 0.991)
    validation        (min: 0.922, max: 0.984, cur: 0.983)
log loss
    training          (min: 0.028, max: 0.764, cur: 0.028)
    validation        (min: 0.058, max: 0.265, cur: 0.064)
Accuracy: 0.9830 (test), 0.9906 (train)

```

For Stochastic Gradient Descent (SGD) and Negative Log-Likelihood Loss Function

In []:

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

model = nn.Sequential()
model.add_module('fc1', nn.Linear(784, 256))
model.add_module('relu1', nn.ReLU())
model.add_module('dropout1', nn.Dropout())
model.add_module('fc2', nn.Linear(256, 256))

```

```

model.add_module('relu2',      nn.ReLU())
model.add_module('dropout2',   nn.Dropout())
model.add_module('fc3',        nn.Linear(256, 10))
print(model)
model.to(device)

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'], flatten=True)
test_set = create_dataset(data['test_x'], data['test_y'], flatten=True)
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.NLLLoss()
optimizer = optim.SGD(model.parameters(), lr=0.001)

liveLoss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

        # Compute the Loss value.
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

        # Update the parameters.
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

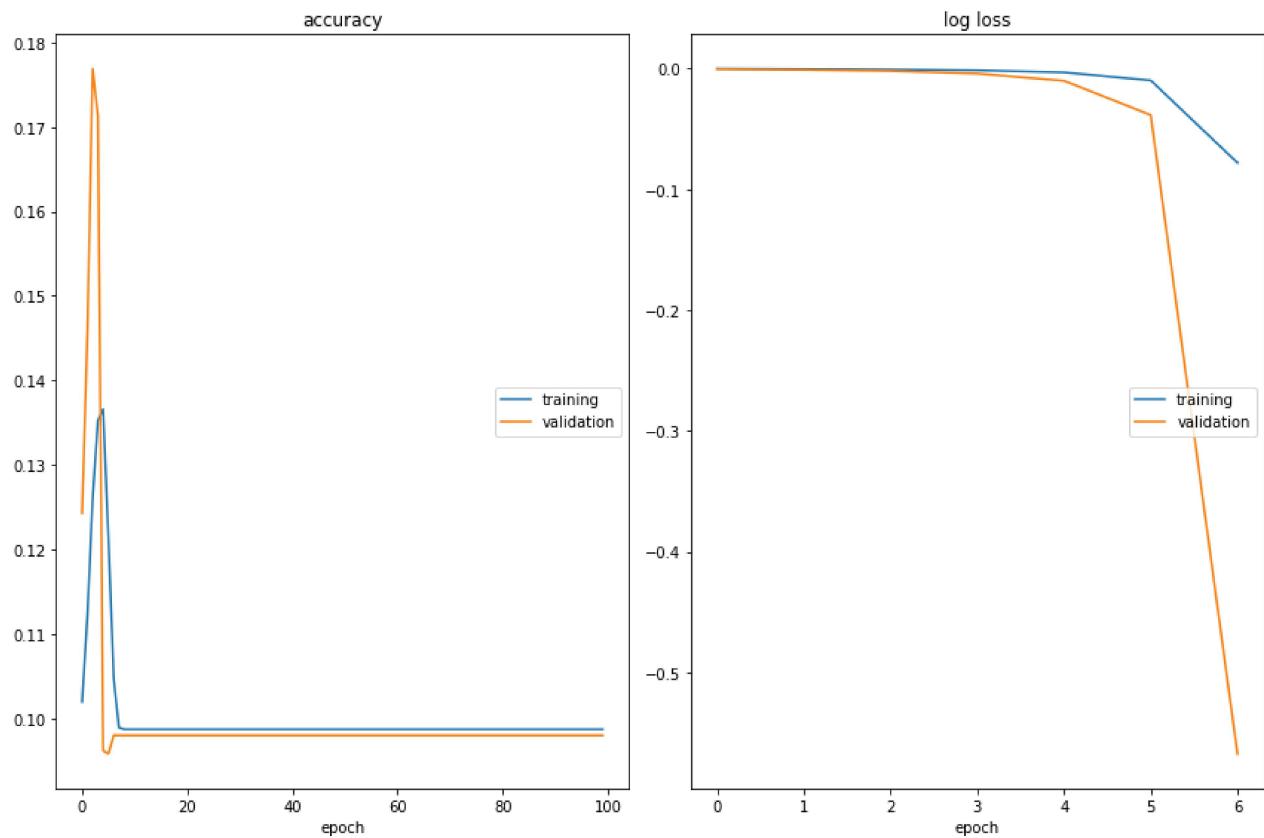
    # Compute the average loss and accuracy.
    train_loss /= len(train_loader.dataset)
    train_correct /= float(len(train_loader.dataset))

    # Evaluate the model on the test set.
    test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

    # Visualize the Loss and accuracy values.
    liveLoss.update({
        'log loss': train_loss,
        'val_log loss': test_loss,
        'accuracy': train_correct,
        'val_accuracy': test_correct,
    })
    liveLoss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))

```



accuracy
 training (min: 0.099, max: 0.137, cur: 0.099)
 validation (min: 0.096, max: 0.177, cur: 0.098)
 log loss
 training (min: -0.078, max: -0.000, cur: nan)
 validation (min: -0.568, max: -0.000, cur: nan)
 Accuracy: 0.0980 (test), 0.0987 (train)

For Adagrad and Multi-Class Cross-Entropy Loss

```
In [ ]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

model = nn.Sequential()
model.add_module('fc1', nn.Linear(784, 256))
model.add_module('relu1', nn.ReLU())
model.add_module('dropout1', nn.Dropout())
model.add_module('fc2', nn.Linear(256, 256))
model.add_module('relu2', nn.ReLU())
model.add_module('dropout2', nn.Dropout())
model.add_module('fc3', nn.Linear(256, 10))
print(model)
model.to(device)
```

```
data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'], flatten=True)
test_set = create_dataset(data['test_x'], data['test_y'], flatten=True)
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.CrossEntropyLoss(size_average=False)
optimizer = optim.Adagrad(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

        # Compute the loss value.
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

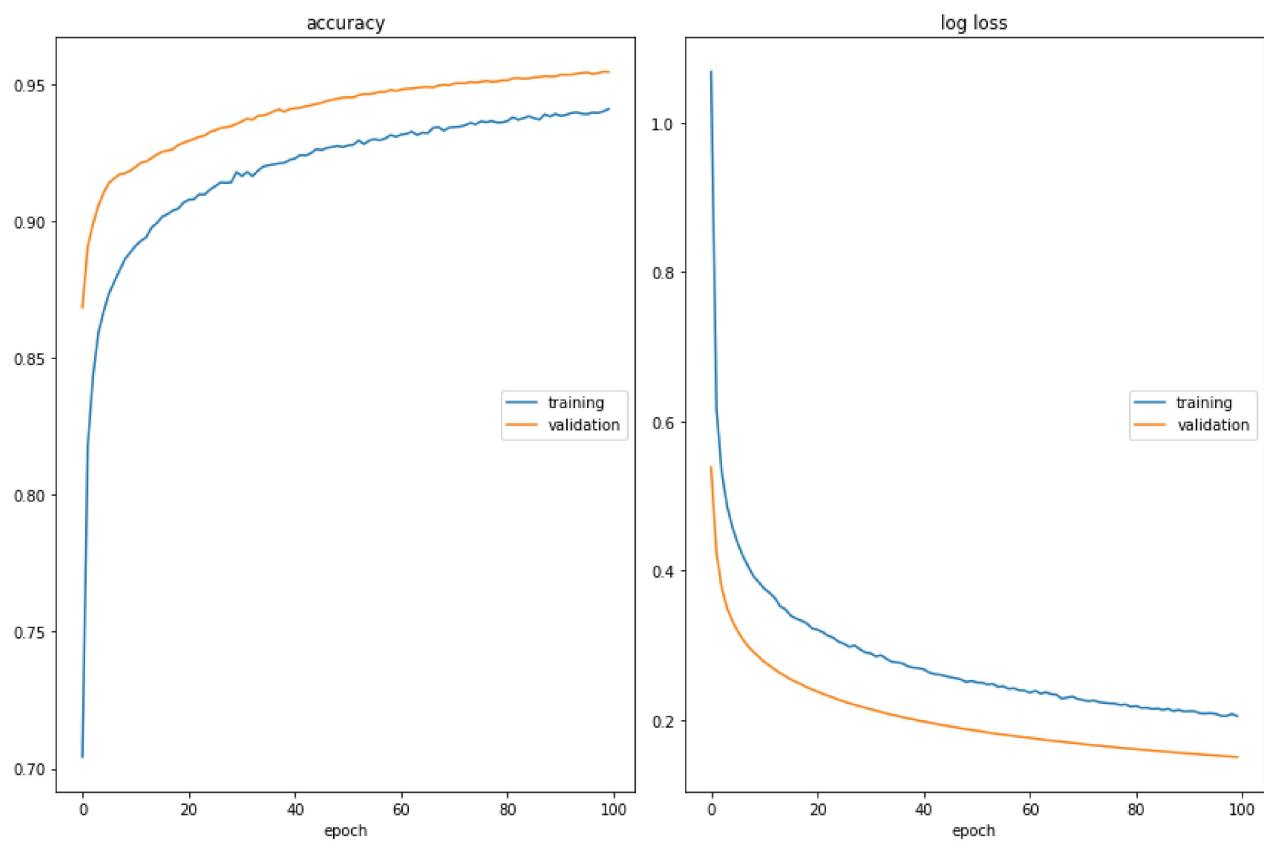
        # Update the parameters.
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    # Compute the average Loss and accuracy.
    train_loss /= len(train_loader.dataset)
    train_correct /= float(len(train_loader.dataset))

    # Evaluate the model on the test set.
    test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

    # Visualize the Loss and accuracy values.
    liveloss.update({
        'log loss': train_loss,
        'val_log loss': test_loss,
        'accuracy': train_correct,
        'val_accuracy': test_correct,
    })
    liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))
```



accuracy
 training (min: 0.704, max: 0.941, cur: 0.941)
 validation (min: 0.869, max: 0.955, cur: 0.954)
 log loss
 training (min: 0.205, max: 1.068, cur: 0.205)
 validation (min: 0.150, max: 0.538, cur: 0.150)
 Accuracy: 0.9544 (test), 0.9408 (train)

For Adagrad and Negative Log-Likelihood Loss Function

In []:

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

model = nn.Sequential()
model.add_module('fc1', nn.Linear(784, 256))
model.add_module('relu1', nn.ReLU())
model.add_module('dropout1', nn.Dropout())
model.add_module('fc2', nn.Linear(256, 256))
model.add_module('relu2', nn.ReLU())
model.add_module('dropout2', nn.Dropout())
model.add_module('fc3', nn.Linear(256, 10))
print(model)
model.to(device)
```

```

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'], flatten=True)
test_set = create_dataset(data['test_x'], data['test_y'], flatten=True)
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.NLLLoss()
optimizer = optim.Adagrad(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

        # Compute the loss value.
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

        # Update the parameters.
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

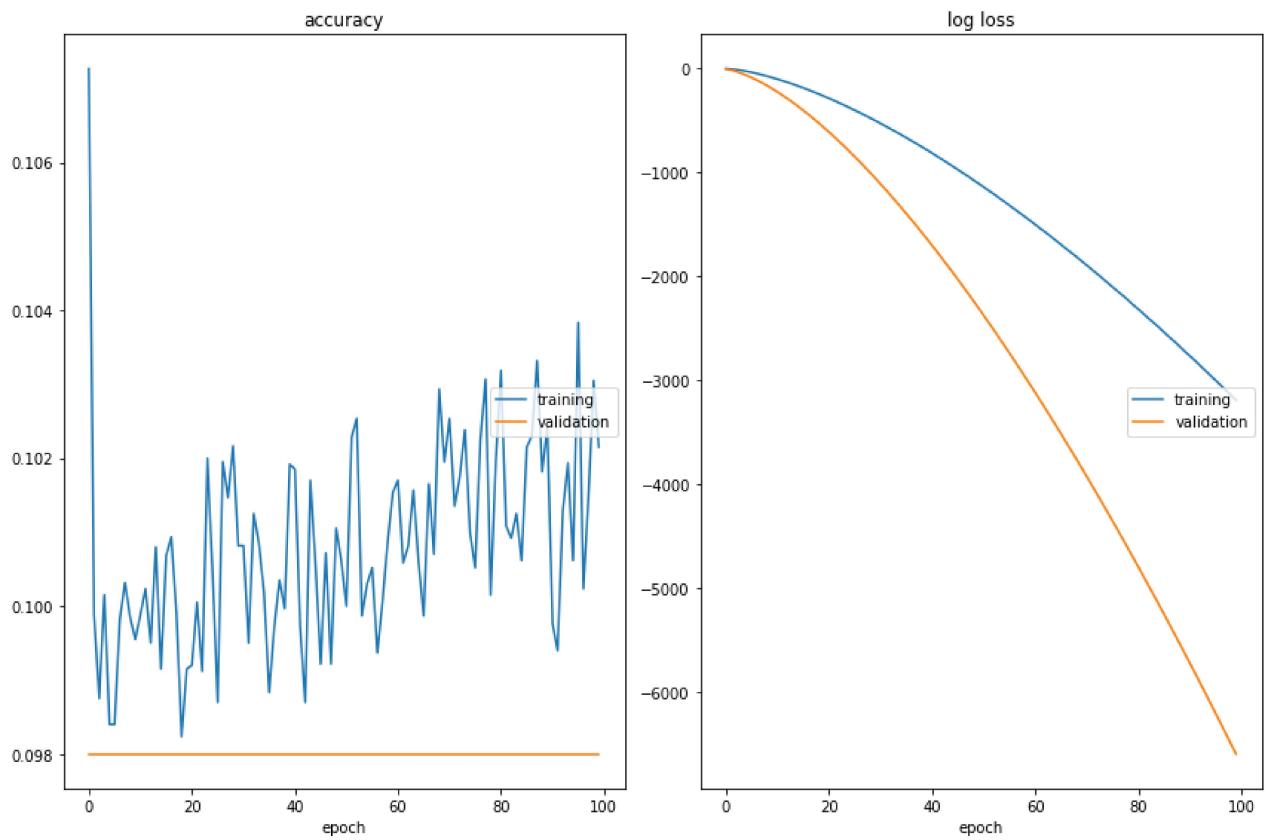
    # Compute the average Loss and accuracy.
    train_loss /= len(train_loader.dataset)
    train_correct /= float(len(train_loader.dataset))

    # Evaluate the model on the test set.
    test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

    # Visualize the Loss and accuracy values.
    liveloss.update({
        'log loss': train_loss,
        'val_log loss': test_loss,
        'accuracy': train_correct,
        'val_accuracy': test_correct,
    })
    liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))

```



```
accuracy
  training          (min: 0.098, max: 0.107, cur: 0.102)
  validation        (min: 0.098, max: 0.098, cur: 0.098)
log loss
  training          (min: -3193.990, max: -0.883, cur: -3193.990)
  validation        (min: -6598.251, max: -4.964, cur: -6598.251)
Accuracy: 0.0980 (test), 0.1022 (train)
```

For RMSProp and Multi-Class Cross-Entropy Loss

In []:

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

model = nn.Sequential()
model.add_module('fc1', nn.Linear(784, 256))
model.add_module('relu1', nn.ReLU())
model.add_module('dropout1', nn.Dropout())
model.add_module('fc2', nn.Linear(256, 256))
model.add_module('relu2', nn.ReLU())
model.add_module('dropout2', nn.Dropout())
model.add_module('fc3', nn.Linear(256, 10))
print(model)
model.to(device)
```

```
data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'], flatten=True)
test_set = create_dataset(data['test_x'], data['test_y'], flatten=True)
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.CrossEntropyLoss(size_average=False)
optimizer = optim.RMSprop(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

        # Compute the loss value.
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

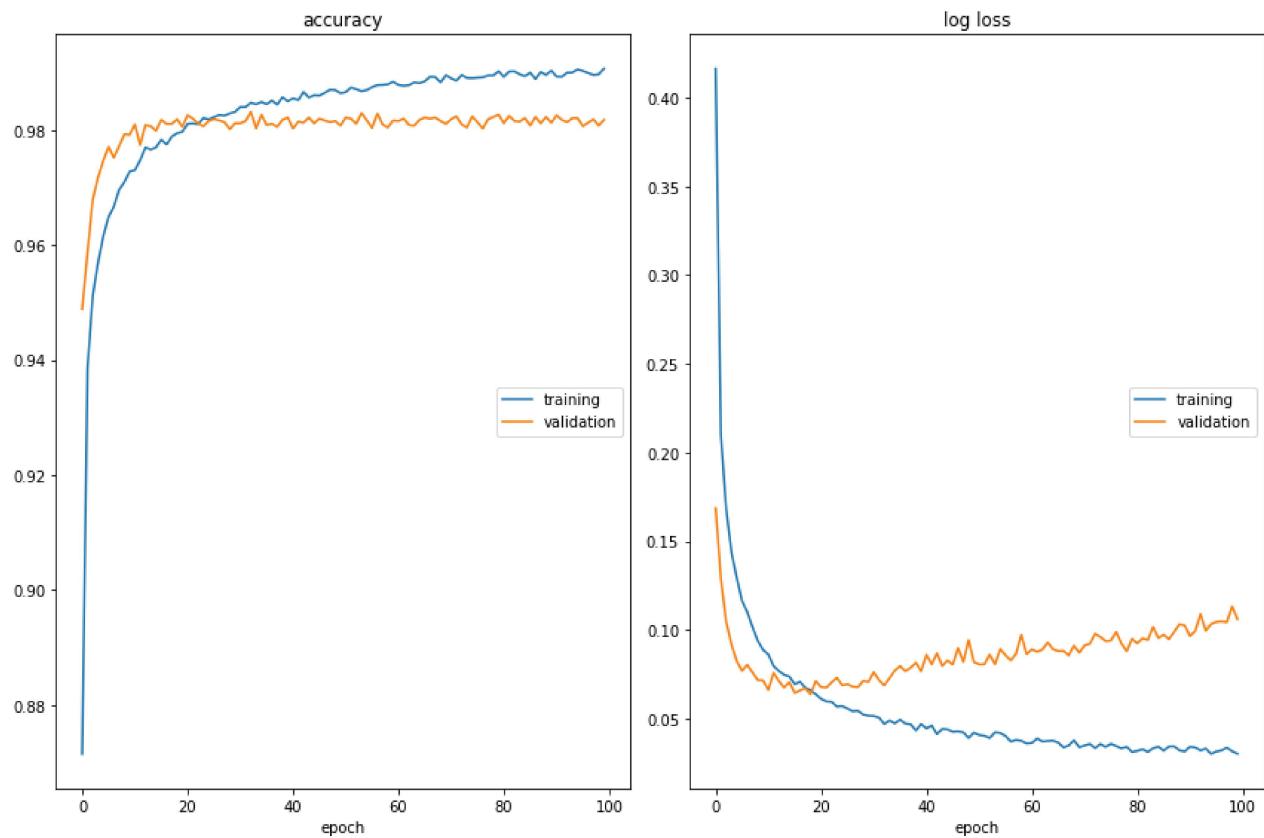
        # Update the parameters.
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    # Compute the average Loss and accuracy.
    train_loss /= len(train_loader.dataset)
    train_correct /= float(len(train_loader.dataset))

    # Evaluate the model on the test set.
    test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

    # Visualize the Loss and accuracy values.
    liveloss.update({
        'log loss': train_loss,
        'val_log loss': test_loss,
        'accuracy': train_correct,
        'val_accuracy': test_correct,
    })
    liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))
```



accuracy
 training (min: 0.871, max: 0.991, cur: 0.991)
 validation (min: 0.949, max: 0.983, cur: 0.982)
 log loss
 training (min: 0.030, max: 0.416, cur: 0.030)
 validation (min: 0.064, max: 0.169, cur: 0.106)
 Accuracy: 0.9818 (test), 0.9907 (train)

For RMSprop and Negative Log-Likelihood Loss Function

In []:

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

model = nn.Sequential()
model.add_module('fc1', nn.Linear(784, 256))
model.add_module('relu1', nn.ReLU())
model.add_module('dropout1', nn.Dropout())
model.add_module('fc2', nn.Linear(256, 256))
model.add_module('relu2', nn.ReLU())
model.add_module('dropout2', nn.Dropout())
model.add_module('fc3', nn.Linear(256, 10))
print(model)
model.to(device)
```

```

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'], flatten=True)
test_set = create_dataset(data['test_x'], data['test_y'], flatten=True)
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.NLLLoss()
optimizer = optim.RMSprop(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

        # Compute the loss value.
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

        # Update the parameters.
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

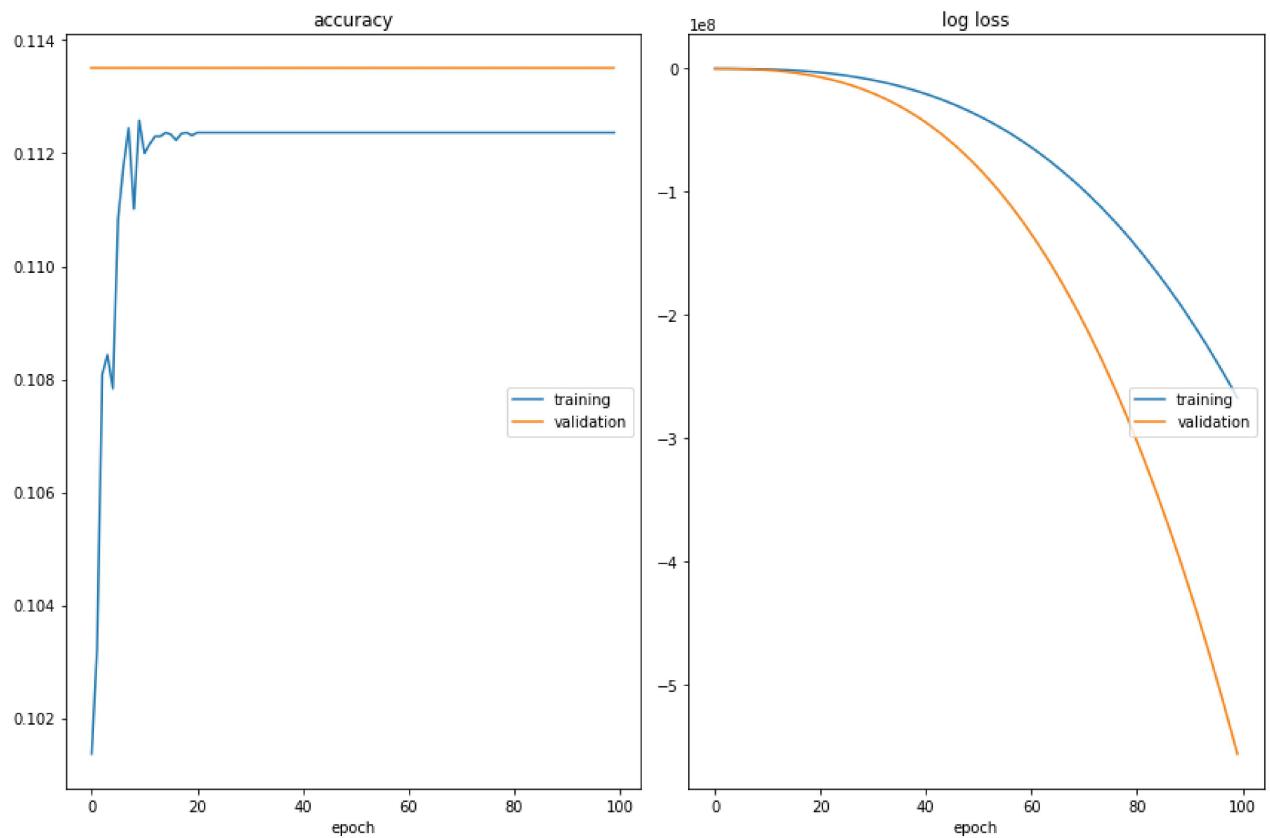
    # Compute the average Loss and accuracy.
    train_loss /= len(train_loader.dataset)
    train_correct /= float(len(train_loader.dataset))

    # Evaluate the model on the test set.
    test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

    # Visualize the Loss and accuracy values.
    liveloss.update({
        'log loss': train_loss,
        'val_log loss': test_loss,
        'accuracy': train_correct,
        'val_accuracy': test_correct,
    })
    liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))

```



```

accuracy
    training          (min: 0.101, max: 0.113, cur: 0.112)
    validation        (min: 0.114, max: 0.114, cur: 0.114)
log loss
    training          (min: -267166984.329, max: -1136.229, cur: -267166984.3
29)
    validation        (min: -555961161.318, max: -6648.717, cur: -555961161.3
18)
Accuracy: 0.1135 (test), 0.1124 (train)

```

For Adam and Multi-Class Cross-Entropy Loss

In []:

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

model = nn.Sequential()
model.add_module('fc1', nn.Linear(784, 256))
model.add_module('relu1', nn.ReLU())
model.add_module('dropout1', nn.Dropout())
model.add_module('fc2', nn.Linear(256, 256))
model.add_module('relu2', nn.ReLU())
model.add_module('dropout2', nn.Dropout())
model.add_module('fc3', nn.Linear(256, 10))
print(model)
model.to(device)

```

```
data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'], flatten=True)
test_set = create_dataset(data['test_x'], data['test_y'], flatten=True)
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.CrossEntropyLoss(size_average=False)
optimizer = optim.Adam(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

        # Compute the loss value.
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

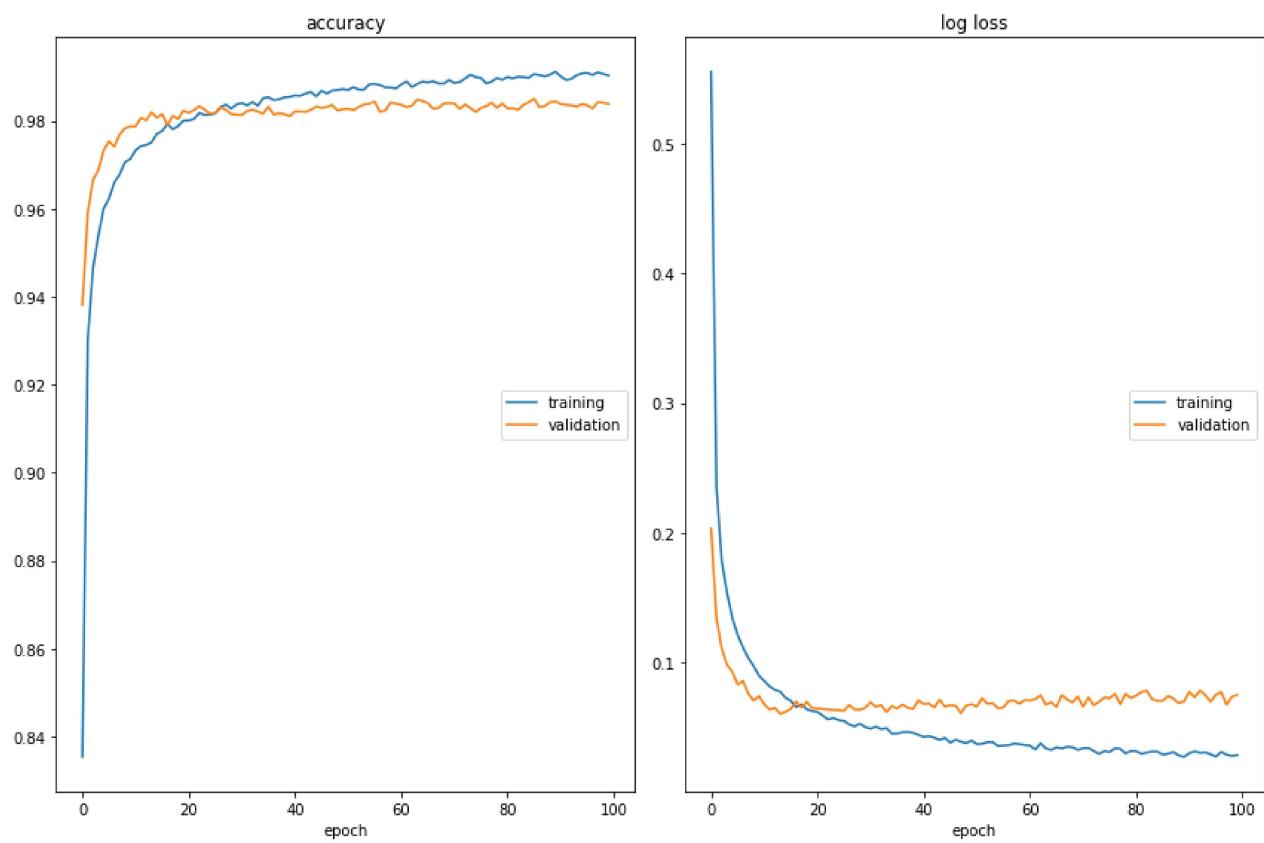
        # Update the parameters.
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    # Compute the average Loss and accuracy.
    train_loss /= len(train_loader.dataset)
    train_correct /= float(len(train_loader.dataset))

    # Evaluate the model on the test set.
    test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

    # Visualize the Loss and accuracy values.
    liveloss.update({
        'log loss': train_loss,
        'val_log loss': test_loss,
        'accuracy': train_correct,
        'val_accuracy': test_correct,
    })
    liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))
```



```
accuracy
  training      (min: 0.835, max: 0.991, cur: 0.990)
  validation    (min: 0.938, max: 0.985, cur: 0.984)
log loss
  training      (min: 0.027, max: 0.555, cur: 0.029)
  validation    (min: 0.061, max: 0.203, cur: 0.075)
Accuracy: 0.9838 (test), 0.9903 (train)
```

For Adam and Negative Log-Likelihood Function

In []:

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

model = nn.Sequential()
model.add_module('fc1', nn.Linear(784, 256))
model.add_module('relu1', nn.ReLU())
model.add_module('dropout1', nn.Dropout())
model.add_module('fc2', nn.Linear(256, 256))
model.add_module('relu2', nn.ReLU())
model.add_module('dropout2', nn.Dropout())
model.add_module('fc3', nn.Linear(256, 10))
print(model)
model.to(device)
```

```
data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'], flatten=True)
test_set = create_dataset(data['test_x'], data['test_y'], flatten=True)
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.NLLLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

        # Compute the loss value.
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

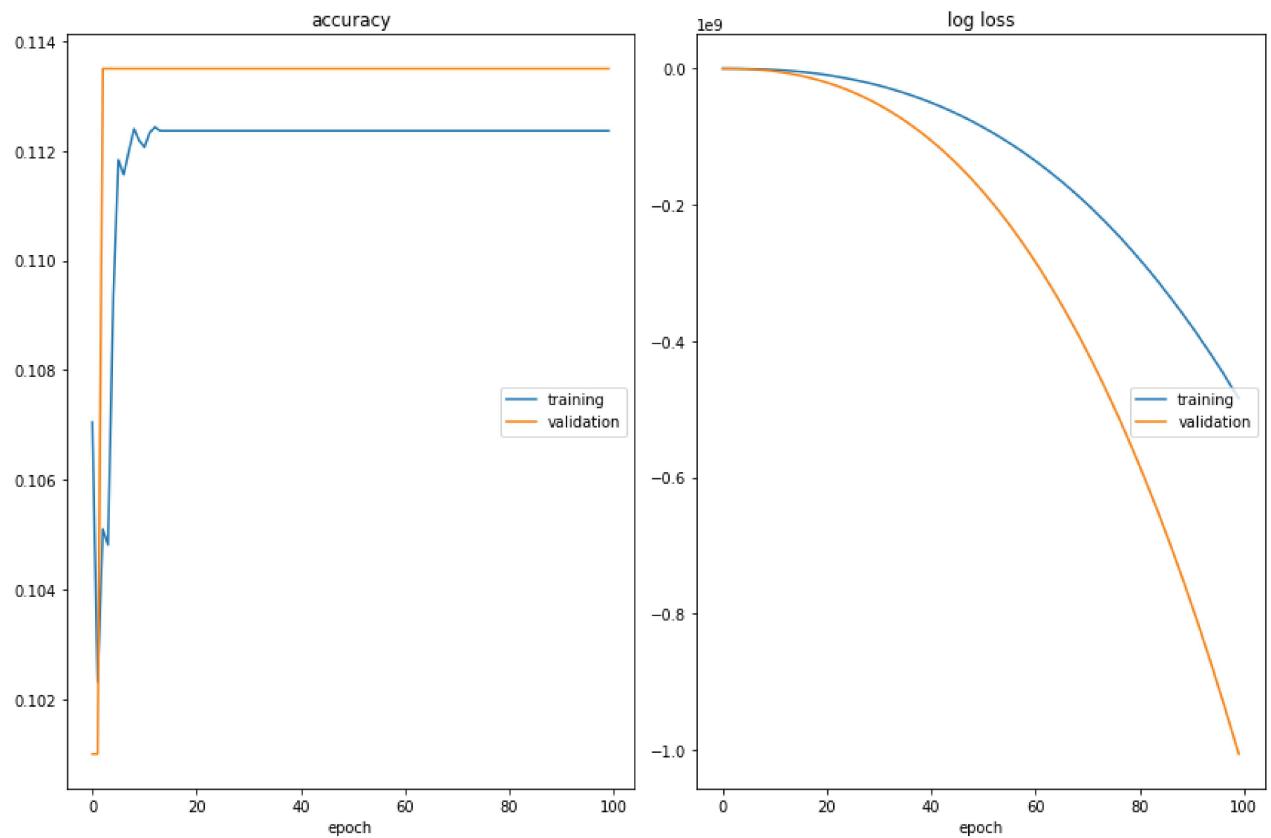
        # Update the parameters.
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    # Compute the average Loss and accuracy.
    train_loss /= len(train_loader.dataset)
    train_correct /= float(len(train_loader.dataset))

    # Evaluate the model on the test set.
    test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

    # Visualize the Loss and accuracy values.
    liveloss.update({
        'log loss': train_loss,
        'val_log loss': test_loss,
        'accuracy': train_correct,
        'val_accuracy': test_correct,
    })
    liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))
```



```

accuracy
  training          (min: 0.102, max: 0.112, cur: 0.112)
  validation        (min: 0.101, max: 0.114, cur: 0.114)
log loss
  training          (min: -483891027.149, max: -307.360, cur: -483891027.14
9)
  validation        (min: -1005867514.266, max: -2963.845, cur: -100586751
4.266)
Accuracy: 0.1135 (test), 0.1124 (train)

```

Comparison table for ANN using different Optimizers and Loss Function

Optimizer	Loss Function	Training Accuracy	Testing Accuracy
Stochastic Gradient Descent (SGD)	Multi-Class Cross-Entropy Loss	0.9906	0.9830
Stochastic Gradient Descent (SGD)	Negative Log-Likelihood Loss Function	0.0987	0.0980
Adagrad	Multi-Class Cross-Entropy Loss	0.9408	0.9544
Adagrad	Negative Log-Likelihood Loss Function	0.1022	0.0980
RMSPROP	Multi-Class Cross-Entropy Loss	0.9907	0.9818
RMSPROP	Negative Log-Likelihood Loss Function	0.1124	0.1135
Adam	Multi-Class Cross-Entropy Loss	0.9903	0.9838
Adam	Negative Log-Likelihood Loss Function	0.1124	0.1135

We got better results by using RMSPROP & Multi-Class Cross-Entropy Loss Function and Adam & Multi-Class Cross-Entropy Loss Function in ANN.

Convolutional Neural Network (CNN)

For Stochastic Gradient Descent (SGD) and Multi-Class Cross-Entropy Loss

In []:

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU
```

```

class Flatten(torch.nn.Module):
    def __init__(self):
        super(Flatten, self).__init__()

    def forward(self, x):
        return x.view(-1, 512)

model = torch.nn.Sequential(
    torch.nn.Conv2d(1, 16, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Conv2d(16, 32, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    Flatten(),
    torch.nn.Linear(512, 256),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Linear(256, 10),
)
print(model)
model.to(device)

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'])
test_set = create_dataset(data['test_x'], data['test_y'])
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.CrossEntropyLoss(size_average=False)
optimizer = optim.SGD(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

        # Compute the loss value.
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

        # Update the parameters.
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    # Compute the average loss and accuracy.
    train_loss /= len(train_loader.dataset)
    train_correct /= float(len(train_loader.dataset))

```

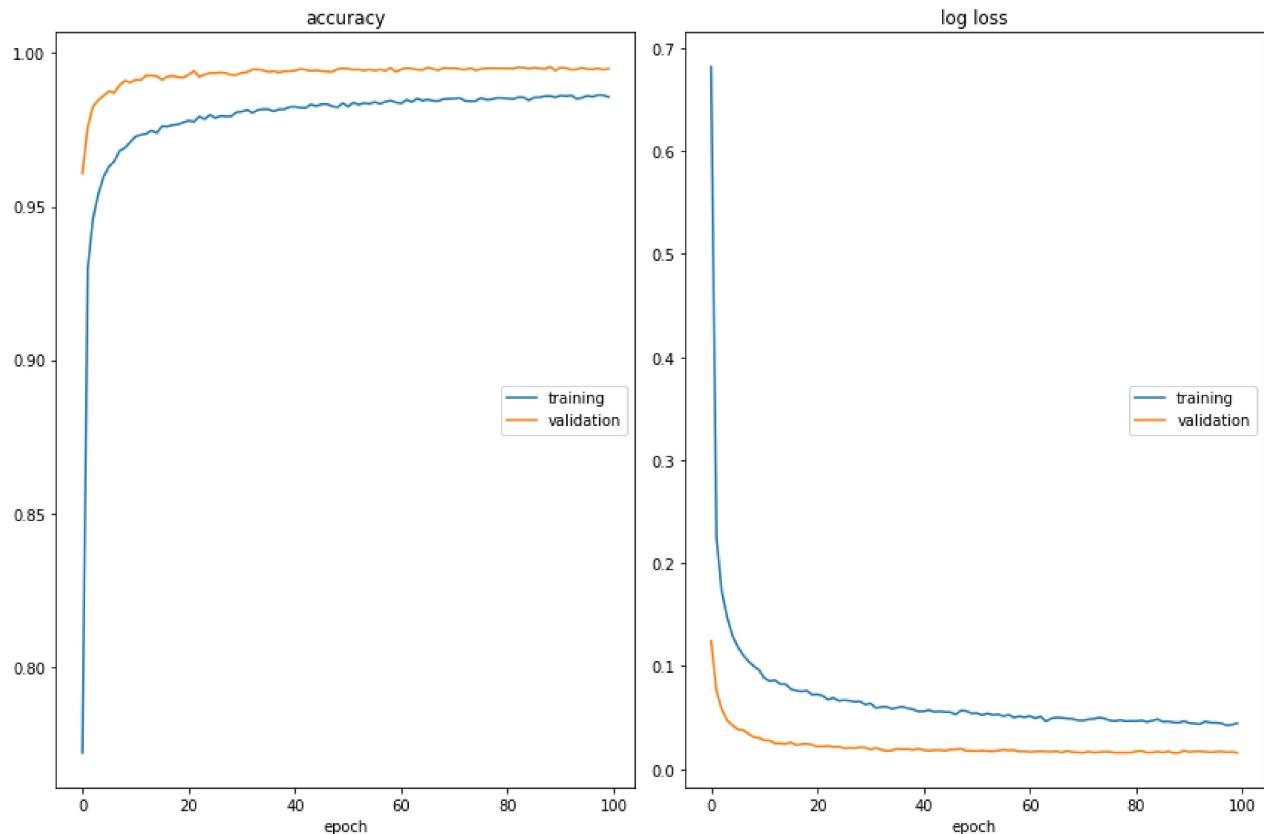
```

# Evaluate the model on the test set.
test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

# Visualize the Loss and accuracy values.
liveloss.update({
    'log loss': train_loss,
    'val_log loss': test_loss,
    'accuracy': train_correct,
    'val_accuracy': test_correct,
})
liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))

```



```

accuracy
    training          (min: 0.772, max: 0.986, cur: 0.986)
    validation        (min: 0.961, max: 0.996, cur: 0.995)
log loss
    training          (min: 0.043, max: 0.682, cur: 0.045)
    validation        (min: 0.016, max: 0.124, cur: 0.016)
Accuracy: 0.9949 (test), 0.9857 (train)

```

For Stochastic Gradient Descent (SGD) and Negative Log-Likelihood Loss Function

In []:

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

```

```

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

class Flatten(torch.nn.Module):
    def __init__(self):
        super(Flatten, self).__init__()

    def forward(self, x):
        return x.view(-1, 512)

model = torch.nn.Sequential(
    torch.nn.Conv2d(1, 16, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Conv2d(16, 32, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    Flatten(),
    torch.nn.Linear(512, 256),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Linear(256, 10),
)
print(model)
model.to(device)

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'])
test_set = create_dataset(data['test_x'], data['test_y'])
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.NLLLoss()
optimizer = optim.SGD(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

        # Compute the loss value.
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

        # Update the parameters.
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    # Compute the average loss and accuracy.

```

```

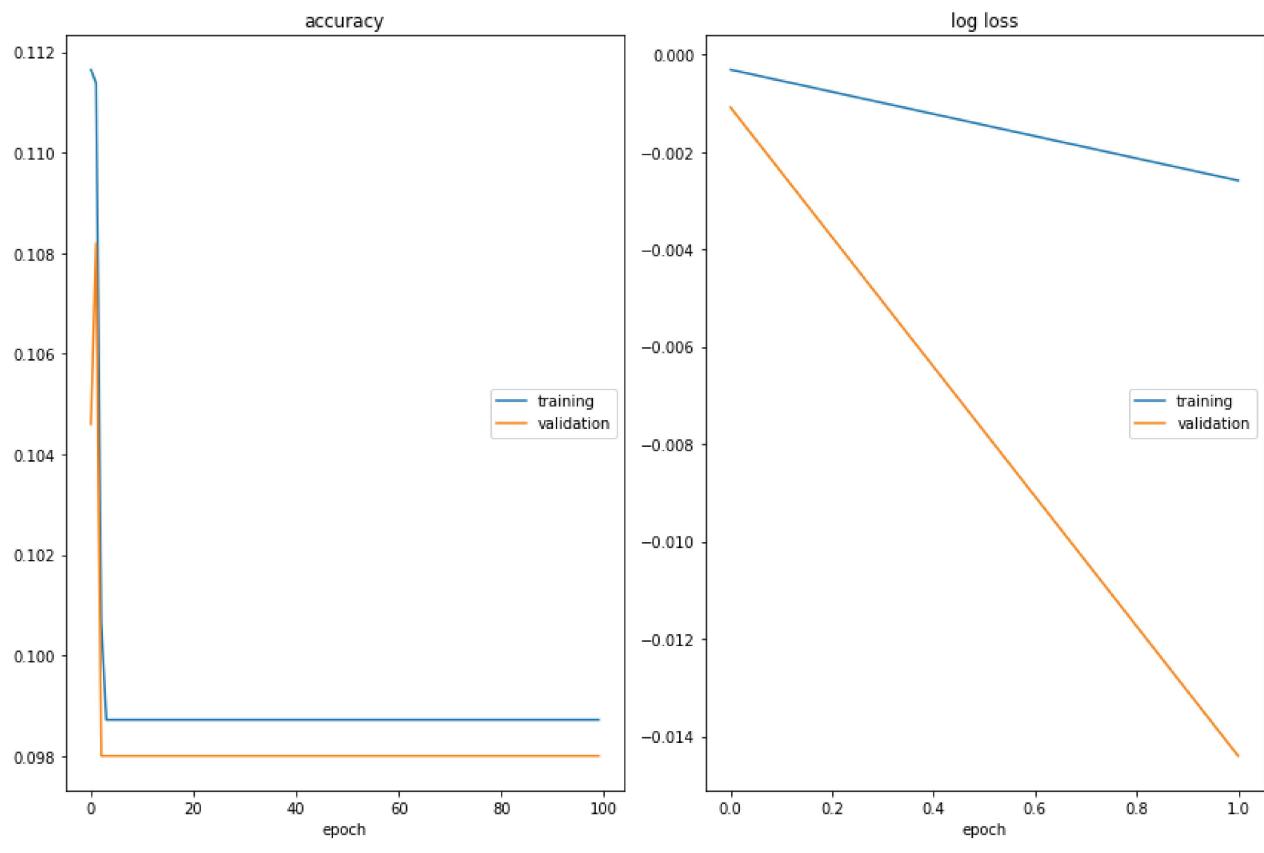
train_loss /= len(train_loader.dataset)
train_correct /= float(len(train_loader.dataset))

# Evaluate the model on the test set.
test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

# Visualize the Loss and accuracy values.
liveloss.update({
    'log loss': train_loss,
    'val_log loss': test_loss,
    'accuracy': train_correct,
    'val_accuracy': test_correct,
})
liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))

```



```

accuracy
    training          (min:  0.099, max:  0.112, cur:  0.099)
    validation        (min:  0.098, max:  0.108, cur:  0.098)
log loss
    training          (min: -0.003, max: -0.000, cur:  nan)
    validation        (min: -0.014, max: -0.001, cur:  nan)
Accuracy: 0.0980 (test), 0.0987 (train)

```

For Adagrad and Multi-Class Cross-Entropy Loss

In []:

```

import torch
import torch.nn as nn
import torch.optim as optim

```

```

from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

class Flatten(torch.nn.Module):
    def __init__(self):
        super(Flatten, self).__init__()

    def forward(self, x):
        return x.view(-1, 512)

model = torch.nn.Sequential(
    torch.nn.Conv2d(1, 16, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Conv2d(16, 32, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    Flatten(),
    torch.nn.Linear(512, 256),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Linear(256, 10),
)
print(model)
model.to(device)

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'])
test_set = create_dataset(data['test_x'], data['test_y'])
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.CrossEntropyLoss(size_average=False)
optimizer = optim.Adagrad(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

        # Compute the Loss value.
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

        # Update the parameters.
        optimizer.zero_grad()
        loss.backward()

```

```

optimizer.step()

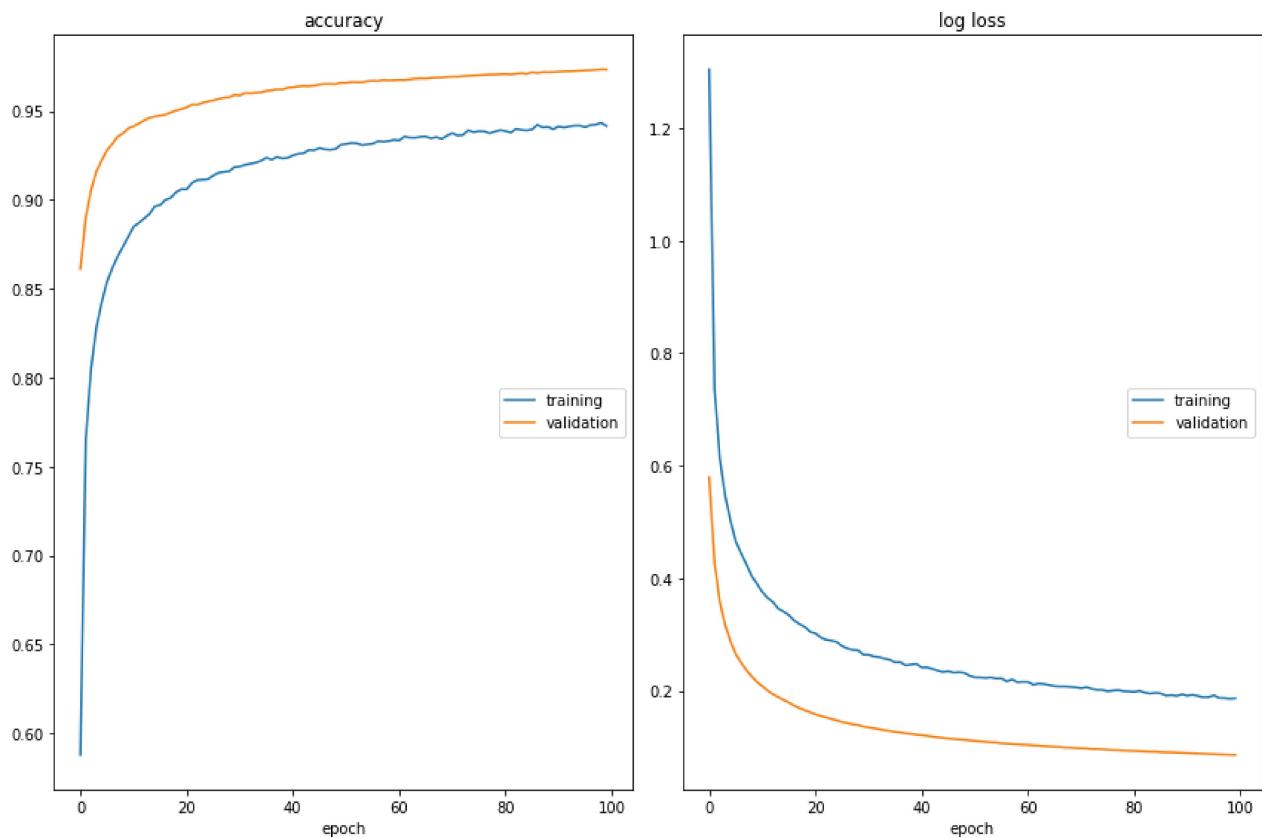
# Compute the average loss and accuracy.
train_loss /= len(train_loader.dataset)
train_correct /= float(len(train_loader.dataset))

# Evaluate the model on the test set.
test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

# Visualize the loss and accuracy values.
liveloss.update({
    'log loss': train_loss,
    'val_log loss': test_loss,
    'accuracy': train_correct,
    'val_accuracy': test_correct,
})
liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))

```



```

accuracy
    training      (min: 0.588, max: 0.943, cur: 0.942)
    validation   (min: 0.861, max: 0.974, cur: 0.973)
log loss
    training      (min: 0.186, max: 1.304, cur: 0.187)
    validation   (min: 0.086, max: 0.579, cur: 0.086)
Accuracy: 0.9734 (test), 0.9416 (train)

```

For Adagrad and Negative Log-Likelihood Loss Function

In []:

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

class Flatten(torch.nn.Module):
    def __init__(self):
        super(Flatten, self).__init__()

    def forward(self, x):
        return x.view(-1, 512)

model = torch.nn.Sequential(
    torch.nn.Conv2d(1, 16, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Conv2d(16, 32, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    Flatten(),
    torch.nn.Linear(512, 256),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Linear(256, 10),
)
print(model)
model.to(device)

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'])
test_set = create_dataset(data['test_x'], data['test_y'])
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.NLLLoss()
optimizer = optim.Adagrad(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

        # Compute the Loss value.
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

```

```

# Update the parameters.
optimizer.zero_grad()
loss.backward()
optimizer.step()

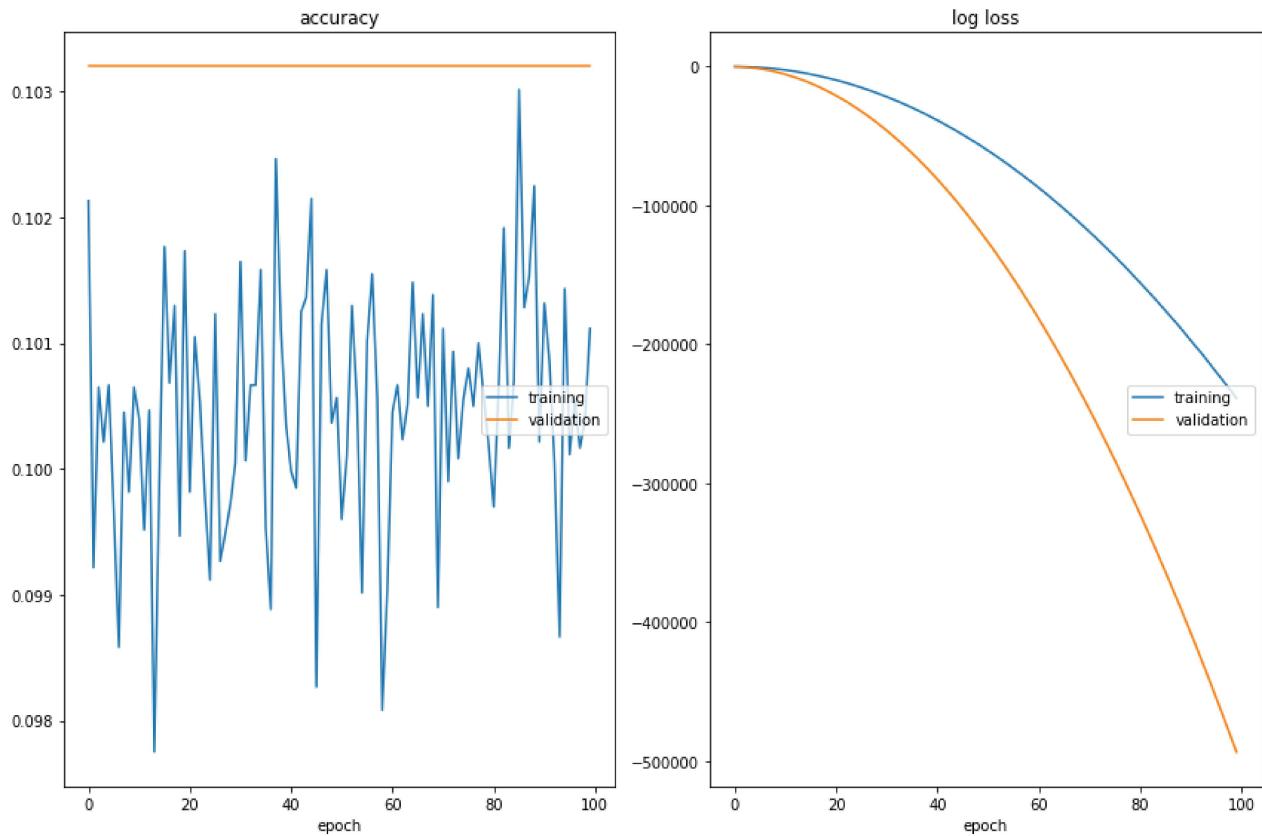
# Compute the average loss and accuracy.
train_loss /= len(train_loader.dataset)
train_correct /= float(len(train_loader.dataset))

# Evaluate the model on the test set.
test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

# Visualize the loss and accuracy values.
liveloss.update({
    'log loss': train_loss,
    'val_log loss': test_loss,
    'accuracy': train_correct,
    'val_accuracy': test_correct,
})
liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))

```



```

accuracy
  training          (min:  0.098, max:  0.103, cur:  0.101)
  validation        (min:  0.103, max:  0.103, cur:  0.103)
log loss
  training          (min: -238921.400, max: -6.397, cur: -238921.400)
  validation        (min: -493492.724, max: -40.992, cur: -493492.724)
Accuracy: 0.1032 (test), 0.1011 (train)

```

For RMSProp and Multi-Class Cross-Entropy

LOSS

In []:

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

class Flatten(torch.nn.Module):
    def __init__(self):
        super(Flatten, self).__init__()

    def forward(self, x):
        return x.view(-1, 512)

model = torch.nn.Sequential(
    torch.nn.Conv2d(1, 16, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Conv2d(16, 32, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    Flatten(),
    torch.nn.Linear(512, 256),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Linear(256, 10),
)
print(model)
model.to(device)

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'])
test_set = create_dataset(data['test_x'], data['test_y'])
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.CrossEntropyLoss(size_average=False)
optimizer = optim.RMSprop(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)
        train_correct += (predicted == y).sum().item()

```

```

# Compute the Loss value.
loss = loss_fn(y_pred, y)
train_loss += loss.item()

# Update the parameters.
optimizer.zero_grad()
loss.backward()
optimizer.step()

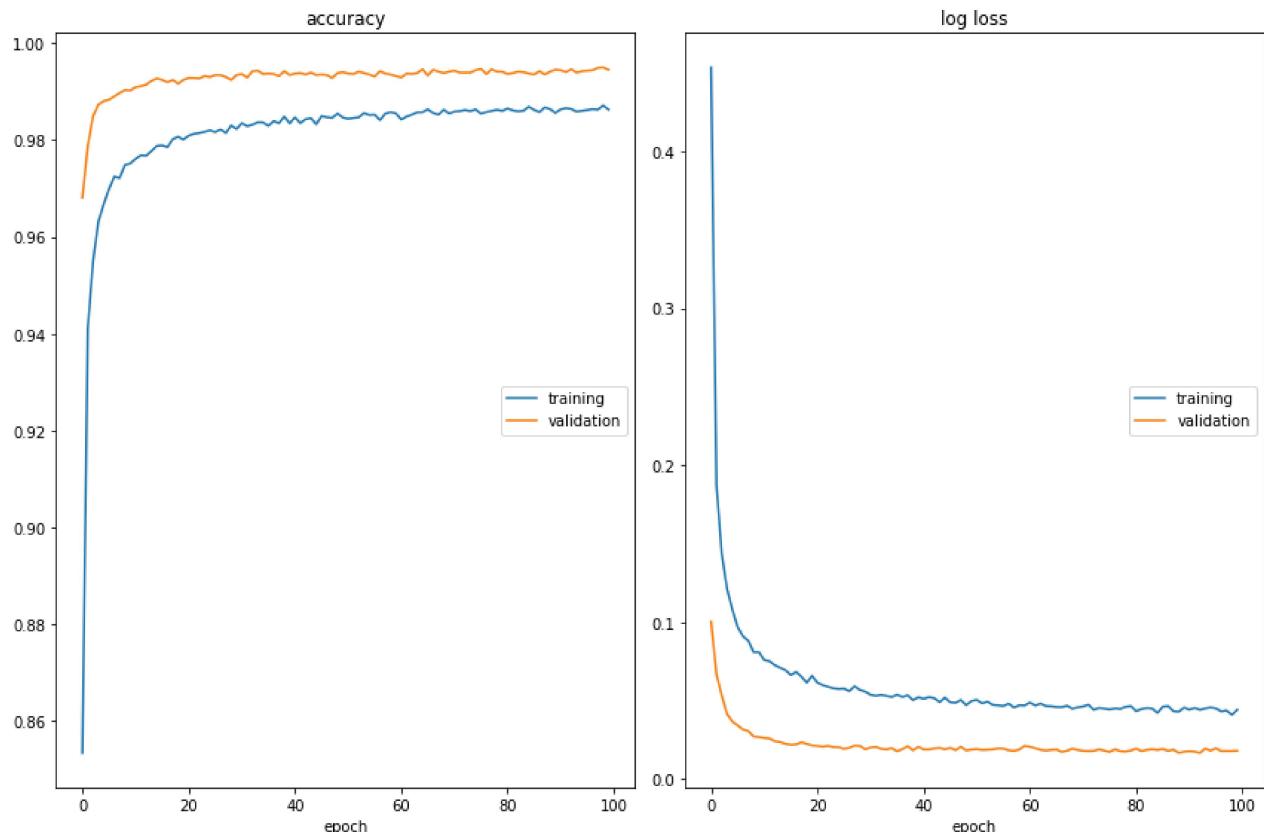
# Compute the average loss and accuracy.
train_loss /= len(train_loader.dataset)
train_correct /= float(len(train_loader.dataset))

# Evaluate the model on the test set.
test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

# Visualize the loss and accuracy values.
liveloss.update({
    'log loss': train_loss,
    'val_log loss': test_loss,
    'accuracy': train_correct,
    'val_accuracy': test_correct,
})
liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))

```



```

accuracy
    training          (min:  0.853, max:  0.987, cur:  0.986)
    validation        (min:  0.968, max:  0.995, cur:  0.995)
log loss
    training          (min:  0.041, max:  0.454, cur:  0.044)
    validation        (min:  0.017, max:  0.100, cur:  0.018)
Accuracy: 0.9945 (test), 0.9863 (train)

```

For RMSprop and Negative Log-Likelihood Loss Function

In []:

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

class Flatten(torch.nn.Module):
    def __init__(self):
        super(Flatten, self).__init__()

    def forward(self, x):
        return x.view(-1, 512)

model = torch.nn.Sequential(
    torch.nn.Conv2d(1, 16, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Conv2d(16, 32, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    Flatten(),
    torch.nn.Linear(512, 256),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Linear(256, 10),
)
print(model)
model.to(device)

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'])
test_set = create_dataset(data['test_x'], data['test_y'])
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.NLLLoss()
optimizer = optim.RMSprop(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training Loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
        x, y = x.to(device), y.to(device)
        y_pred = model(x)
        _, predicted = torch.max(y_pred.data, 1)

```

```

train_correct += (predicted == y).sum().item()

# Compute the loss value.
loss = loss_fn(y_pred, y)
train_loss += loss.item()

# Update the parameters.
optimizer.zero_grad()
loss.backward()
optimizer.step()

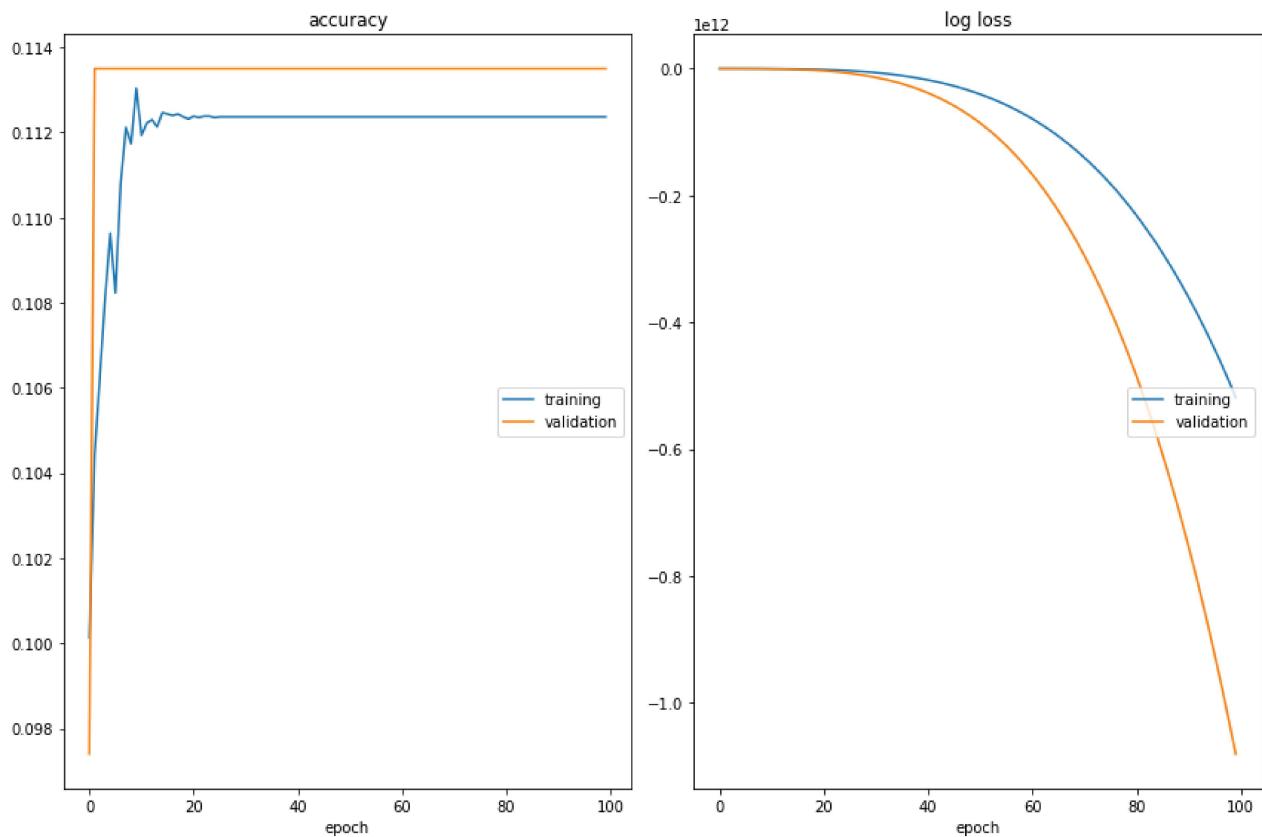
# Compute the average loss and accuracy.
train_loss /= len(train_loader.dataset)
train_correct /= float(len(train_loader.dataset))

# Evaluate the model on the test set.
test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

# Visualize the loss and accuracy values.
liveloss.update({
    'log loss': train_loss,
    'val_log loss': test_loss,
    'accuracy': train_correct,
    'val_accuracy': test_correct,
})
liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))

```



accuracy	
training	(min: 0.100, max: 0.113, cur: 0.112)
validation	(min: 0.097, max: 0.114, cur: 0.114)
log loss	
training	(min: -518730874215.902, max: -68505.597, cur: -5187308)

```
74215.902)
    validation          (min: -1080895272386.560, max: -480433.920, cur: -10808
95272386.560)
Accuracy: 0.1135 (test), 0.1124 (train)
```

For Adam and Multi-Class Cross-Entropy Loss

In []:

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

class Flatten(torch.nn.Module):
    def __init__(self):
        super(Flatten, self).__init__()

    def forward(self, x):
        return x.view(-1, 512)

model = torch.nn.Sequential(
    torch.nn.Conv2d(1, 16, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Conv2d(16, 32, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    Flatten(),
    torch.nn.Linear(512, 256),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Linear(256, 10),
)
print(model)
model.to(device)

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'])
test_set = create_dataset(data['test_x'], data['test_y'])
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.CrossEntropyLoss(size_average=False)
optimizer = optim.Adam(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

    # Training loop for mini-batches
    for batch_idx, (x, y) in enumerate(train_loader):
        # Make predictions with the current parameters.
```

```

x, y = x.to(device), y.to(device)
y_pred = model(x)
_, predicted = torch.max(y_pred.data, 1)
train_correct += (predicted == y).sum().item()

# Compute the loss value.
loss = loss_fn(y_pred, y)
train_loss += loss.item()

# Update the parameters.
optimizer.zero_grad()
loss.backward()
optimizer.step()

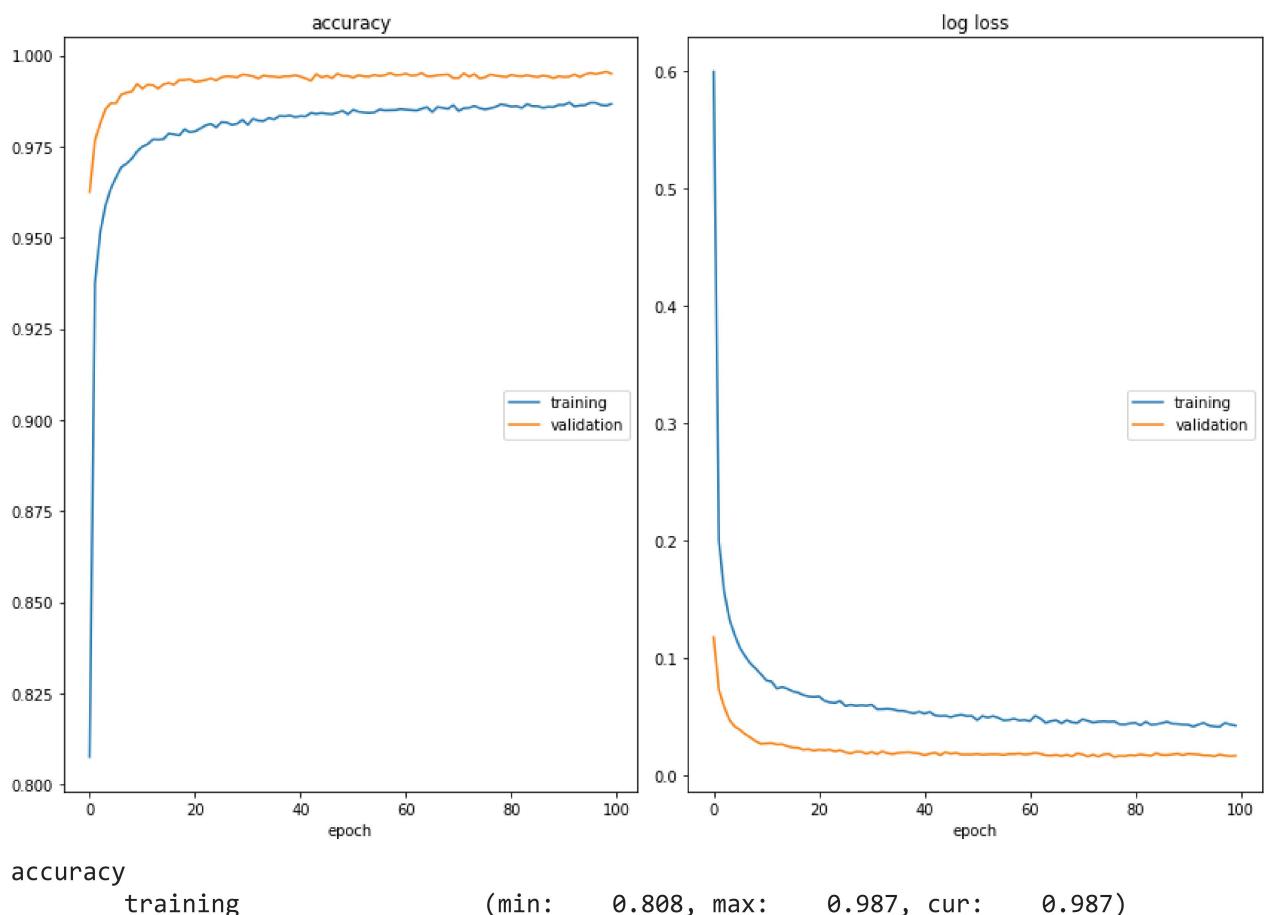
# Compute the average loss and accuracy.
train_loss /= len(train_loader.dataset)
train_correct /= float(len(train_loader.dataset))

# Evaluate the model on the test set.
test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

# Visualize the Loss and accuracy values.
liveloss.update({
    'log loss': train_loss,
    'val_log loss': test_loss,
    'accuracy': train_correct,
    'val_accuracy': test_correct,
})
liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))

```



validation	(min: 0.963, max: 0.996, cur: 0.995)
log loss	
training	(min: 0.041, max: 0.600, cur: 0.042)
validation	(min: 0.015, max: 0.117, cur: 0.016)
Accuracy: 0.9950 (test), 0.9867 (train)	

For Adam and Negative Log-Likelihood Loss Function

In []:

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
import numpy as np
from livelossplot import PlotLosses

device = torch.device("cpu") # Uncomment this to run on CPU
#device = torch.device("cuda:0") # Uncomment this to run on GPU

class Flatten(torch.nn.Module):
    def __init__(self):
        super(Flatten, self).__init__()

    def forward(self, x):
        return x.view(-1, 512)

model = torch.nn.Sequential(
    torch.nn.Conv2d(1, 16, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Conv2d(16, 32, (5, 5)),
    torch.nn.MaxPool2d(2),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    Flatten(),
    torch.nn.Linear(512, 256),
    torch.nn.ReLU(),
    torch.nn.Dropout(),
    torch.nn.Linear(256, 10),
)
print(model)
model.to(device)

data = np.load('mnist.npz')
train_set = create_dataset(data['train_x'], data['train_y'])
test_set = create_dataset(data['test_x'], data['test_y'])
train_loader = DataLoader(train_set, batch_size=256, shuffle=True)
test_loader = DataLoader(test_set, batch_size=128)

loss_fn = nn.NLLLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

liveloss = PlotLosses()
for t in range(100):
    train_loss = 0.
    train_correct = 0

```

```
# Training Loop for mini-batches
for batch_idx, (x, y) in enumerate(train_loader):
    # Make predictions with the current parameters.
    x, y = x.to(device), y.to(device)
    y_pred = model(x)
    _, predicted = torch.max(y_pred.data, 1)
    train_correct += (predicted == y).sum().item()

    # Compute the loss value.
    loss = loss_fn(y_pred, y)
    train_loss += loss.item()

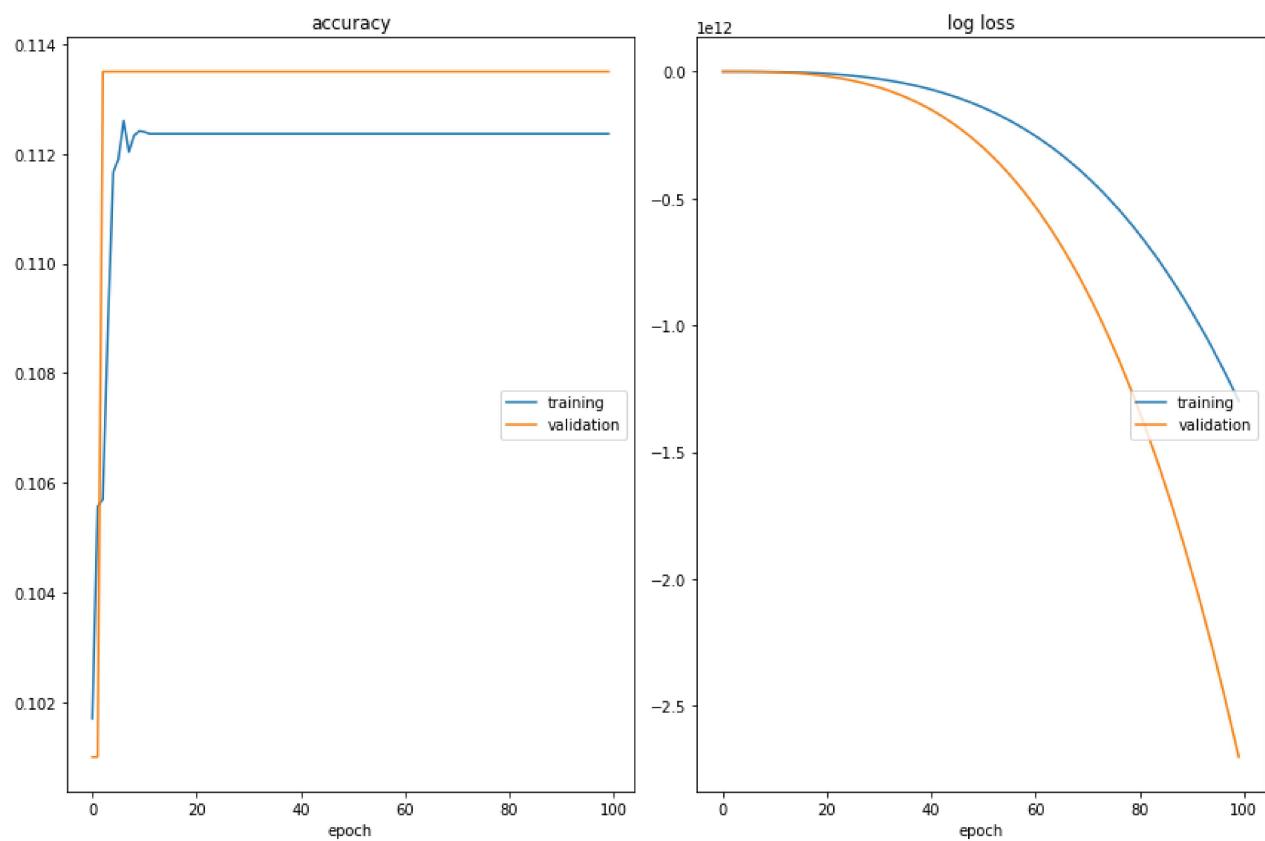
    # Update the parameters.
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

# Compute the average Loss and accuracy.
train_loss /= len(train_loader.dataset)
train_correct /= float(len(train_loader.dataset))

# Evaluate the model on the test set.
test_loss, test_correct = test_model(model, loss_fn, test_loader, device)

# Visualize the Loss and accuracy values.
liveloss.update({
    'log loss': train_loss,
    'val_log loss': test_loss,
    'accuracy': train_correct,
    'val_accuracy': test_correct,
})
liveloss.draw()

print('Accuracy: {:.4f} (test), {:.4f} (train)'.format(test_correct, train_correct))
```



```

accuracy
  training          (min: 0.102, max: 0.113, cur: 0.112)
  validation        (min: 0.101, max: 0.114, cur: 0.114)

log loss
  training          (min: -1298036947034.658, max: -10999.907, cur: -129803
6947034.658)
  validation        (min: -2701167147404.493, max: -136078.734, cur: -27011
67147404.493)

Accuracy: 0.1135 (test), 0.1124 (train)

```

In []:

Comparison table for ANN using different Optimizers and Loss Function

Optimizer	Loss Function	Training Accuracy	Testing Accuracy
Stochastic Gradient Descent (SGD)	Multi-Class Cross-Entropy Loss	0.1124	0.1135
Stochastic Gradient Descent (SGD)	Negative Log-Likelihood Loss Function	0.0987	0.0980
Adagrad	Multi-Class Cross-Entropy Loss	0.9416	0.9734
Adagrad	Negative Log-Likelihood Loss Function	0.1011	0.1032
RMSPROP	Multi-Class Cross-Entropy Loss	0.9863	0.9945
RMSPROP	Negative Log-Likelihood Loss Function	0.1124	0.1135
Adam	Multi-Class Cross-Entropy Loss	0.9867	0.9950
Adam	Negative Log-Likelihood Loss Function	0.1124	0.1135

We got better results by using RMSPROP & Multi-Class Cross-Entropy Loss Function and Adam & Multi-Class Cross-Entropy Loss Function in CNN.